


Chapter 3: Processes





Chapter 3: Processes

- ***Process Concept***
- ***Process Scheduling***
- ***Operations on Processes***
- ***Interprocess Communication***
- ***Examples of IPC Systems***
- ***Communication in Client-Server Systems***






Objectives

- ***To introduce the notion of a process -- a program in execution, which forms the basis of all computation***
- ***To describe the various features of processes, including scheduling, creation and termination, and communication***
- ***To describe communication in client-server systems***





Process Concept

■ ***An operating system executes a variety of programs:***

- Batch system – jobs
- Time-shared systems – user programs or tasks

■ ***Textbook uses the terms job and process almost interchangeably***

■ ***Process - a program in execution; process execution must progress in sequential fashion***

■ ***A process includes:***

- program counter
- stack
- data section





The Process

■ ***Multiple parts***

- The program code, also called **text section**
- Current activity including **program counter**, processor registers
- **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
- **Data section** containing global variables
- **Heap** containing memory dynamically allocated during run time

■ ***Program is passive entity, process is active***

- Program becomes process when executable file loaded into memory

■ ***Execution of program started via GUI mouse clicks, command line entry of its name, etc***


■ ***One program can be several processes***


- Consider multiple users executing the same program





Process in Memory





Process State

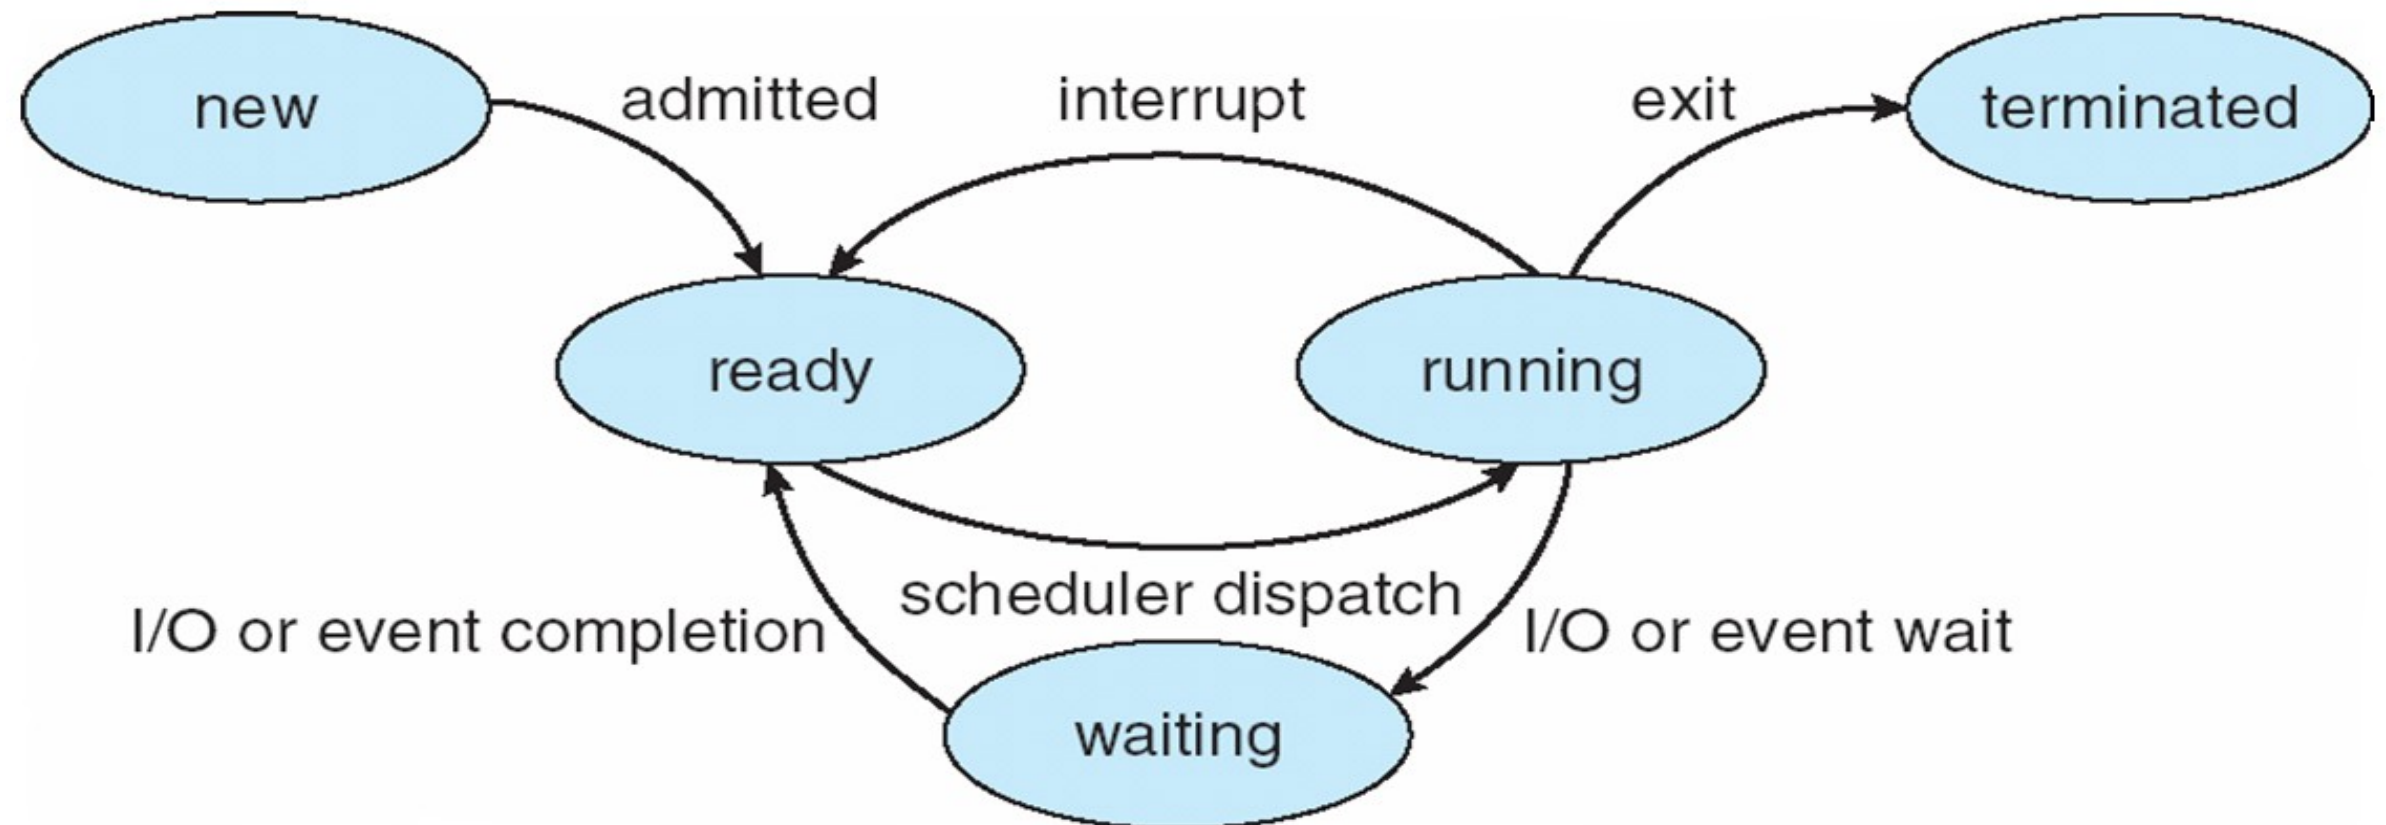
■ ***As a process executes, it changes state***

- **new:** The process is being created
- **running:** Instructions are being executed
- **waiting:** The process is waiting for some event to occur
- **ready:** The process is waiting to be assigned to a processor
- **terminated:** The process has finished execution





Diagram of Process State





Process Control Block (PCB)

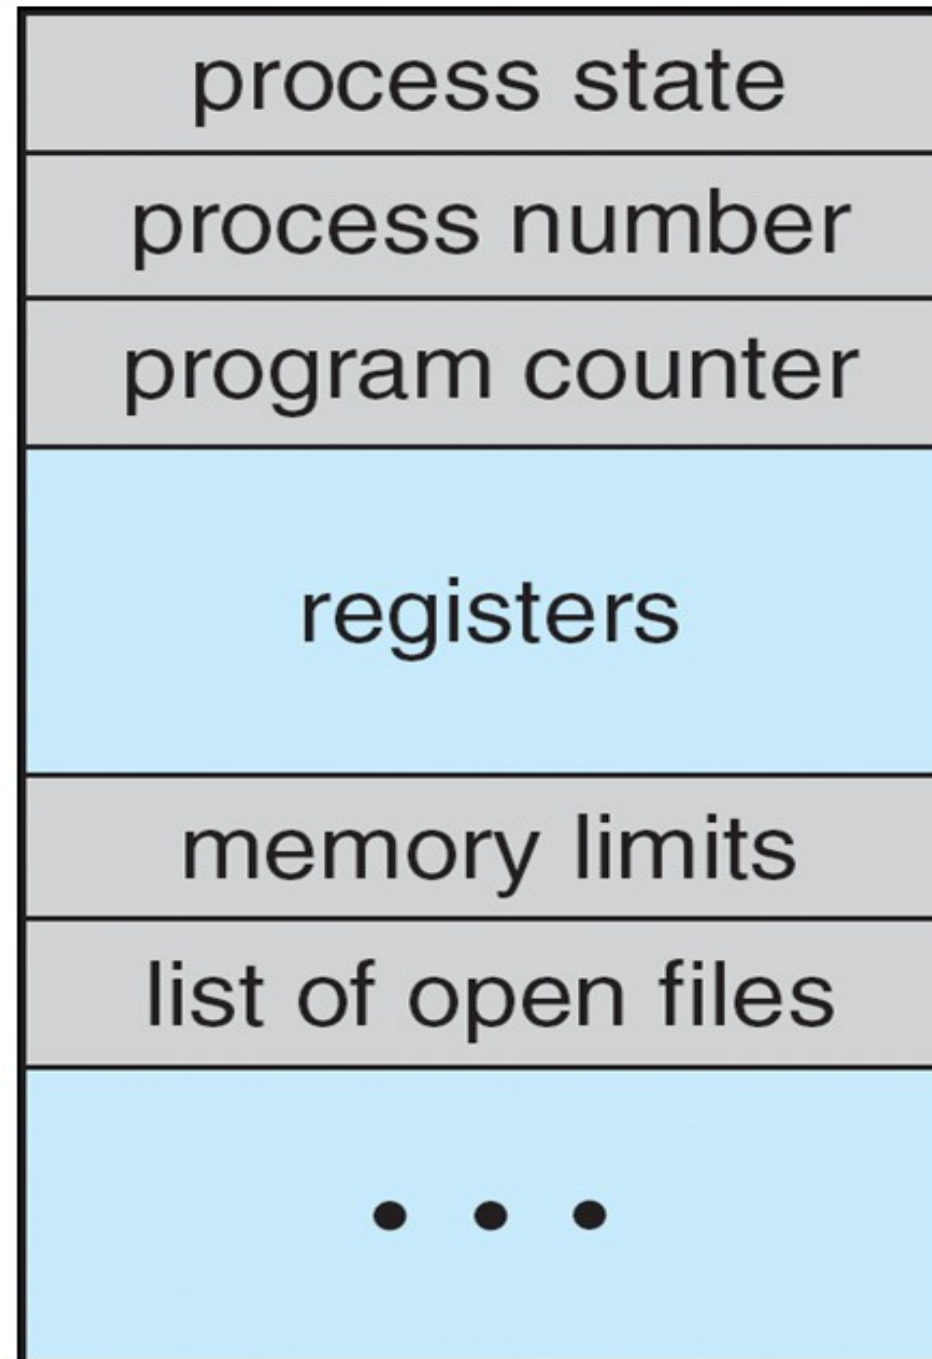
Information associated with each process

- ***Process state***
- ***Program counter***
- ***CPU registers***
- ***CPU scheduling information***
- ***Memory-management information***
- ***Accounting information***
- ***I/O status information***







Process Control Block (PCB)





CPU Switch From Process to Process





Process Scheduling

- ***Maximize CPU use, quickly switch processes onto CPU for time sharing***
- ***Process scheduler selects among available processes for next execution on CPU***
- ***Maintains scheduling queues of processes***
 - Job queue – set of all processes in the system
 - Ready queue – set of all processes residing in main memory, ready and waiting to execute
 - Device queues – set of processes waiting for an I/O device
 - Processes migrate among the various queues

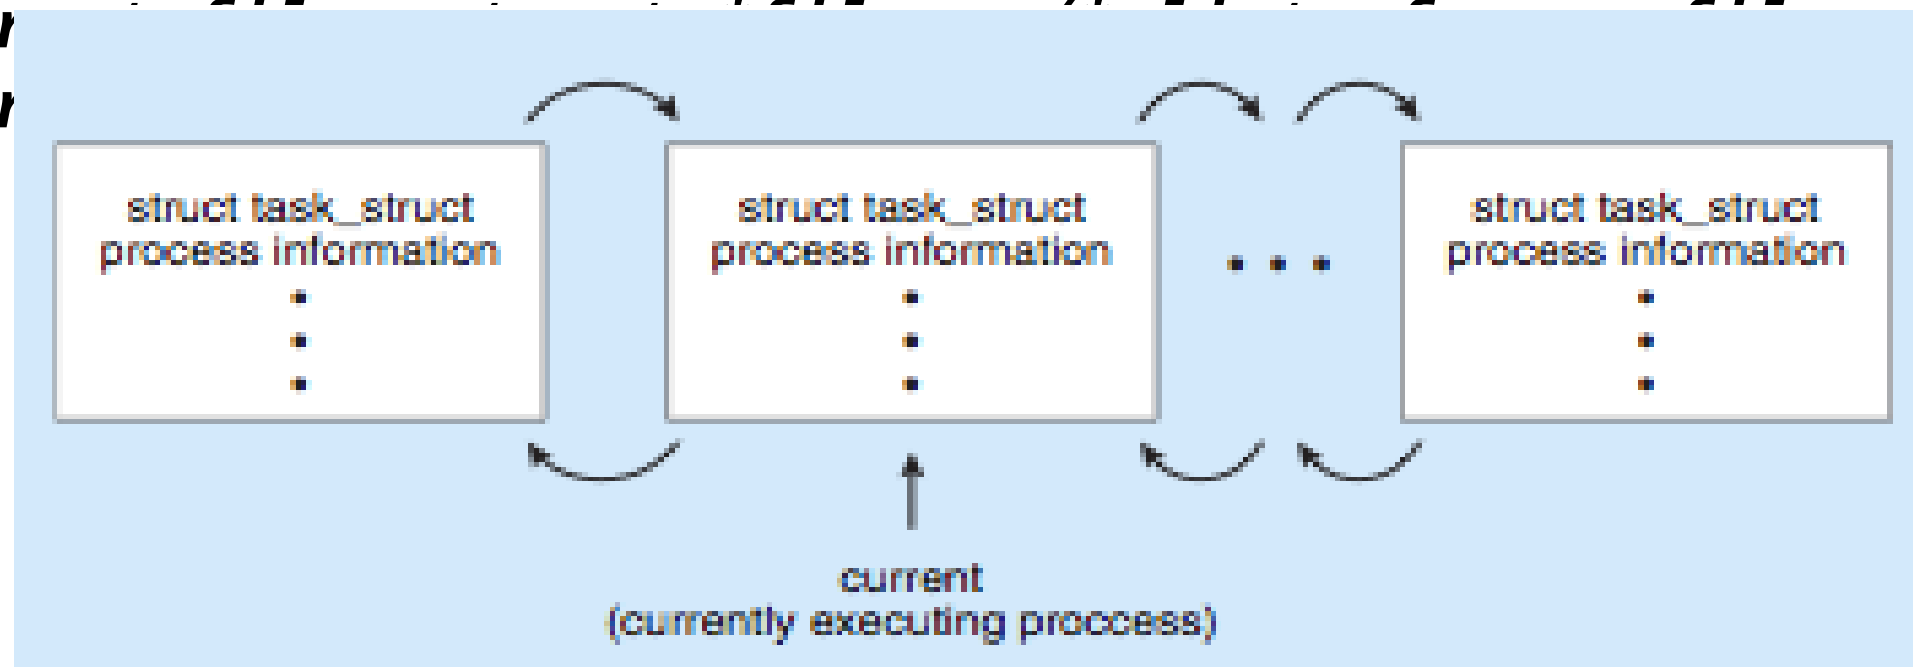




Process Representation in Linux

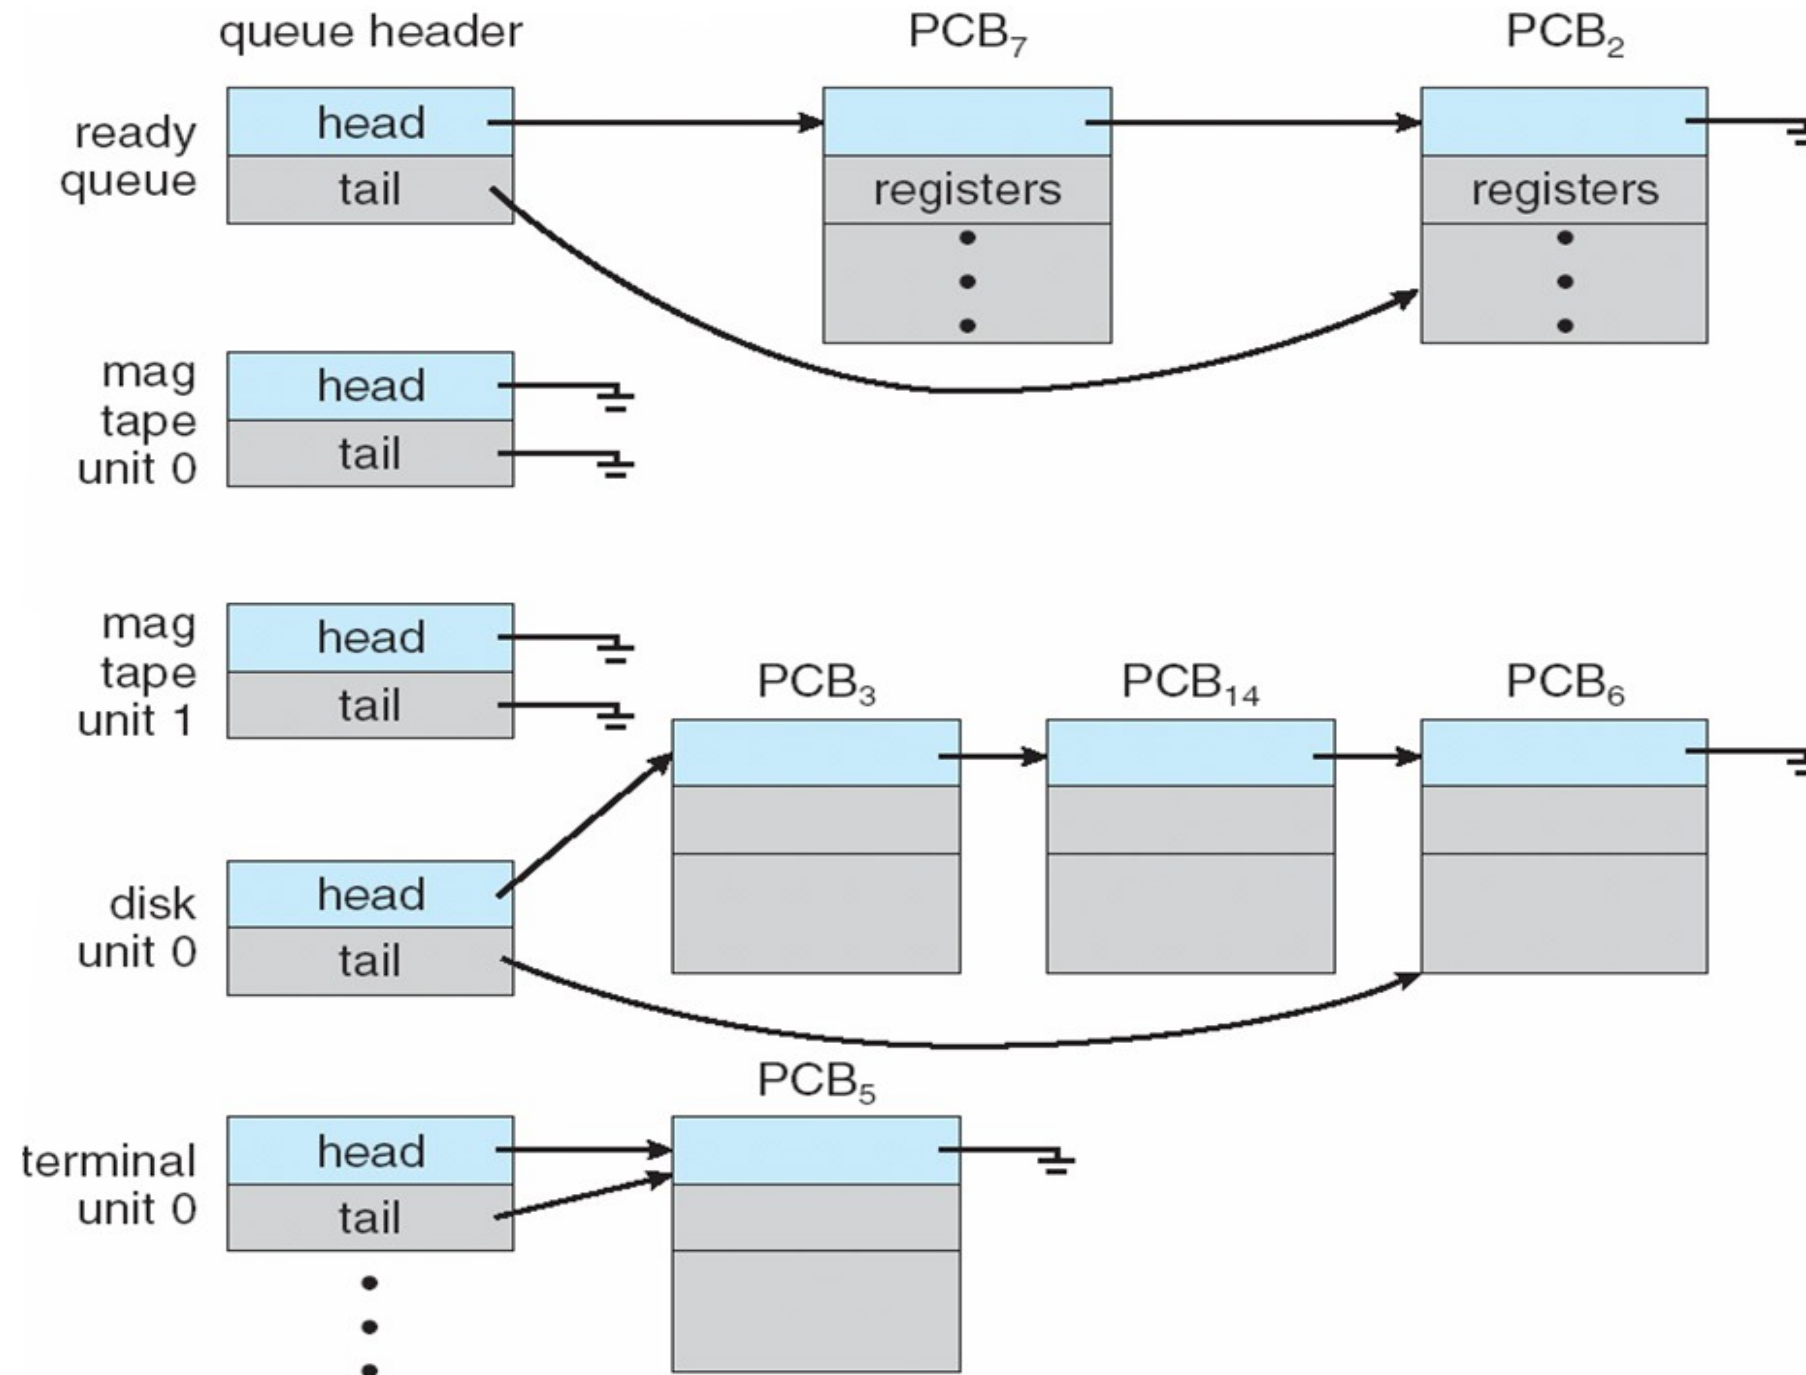
■ Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct ... /* ... */
struct ... /* ... */
```






Ready Queue And Various I/O Device Queues





Representation of Process Scheduling





Schedulers

- ***Long-term scheduler (or job scheduler) - selects which processes should be brought into the ready queue***
- ***Short-term scheduler (or CPU scheduler) - selects which process should be executed next and allocates CPU***
 - Sometimes the only scheduler in a system





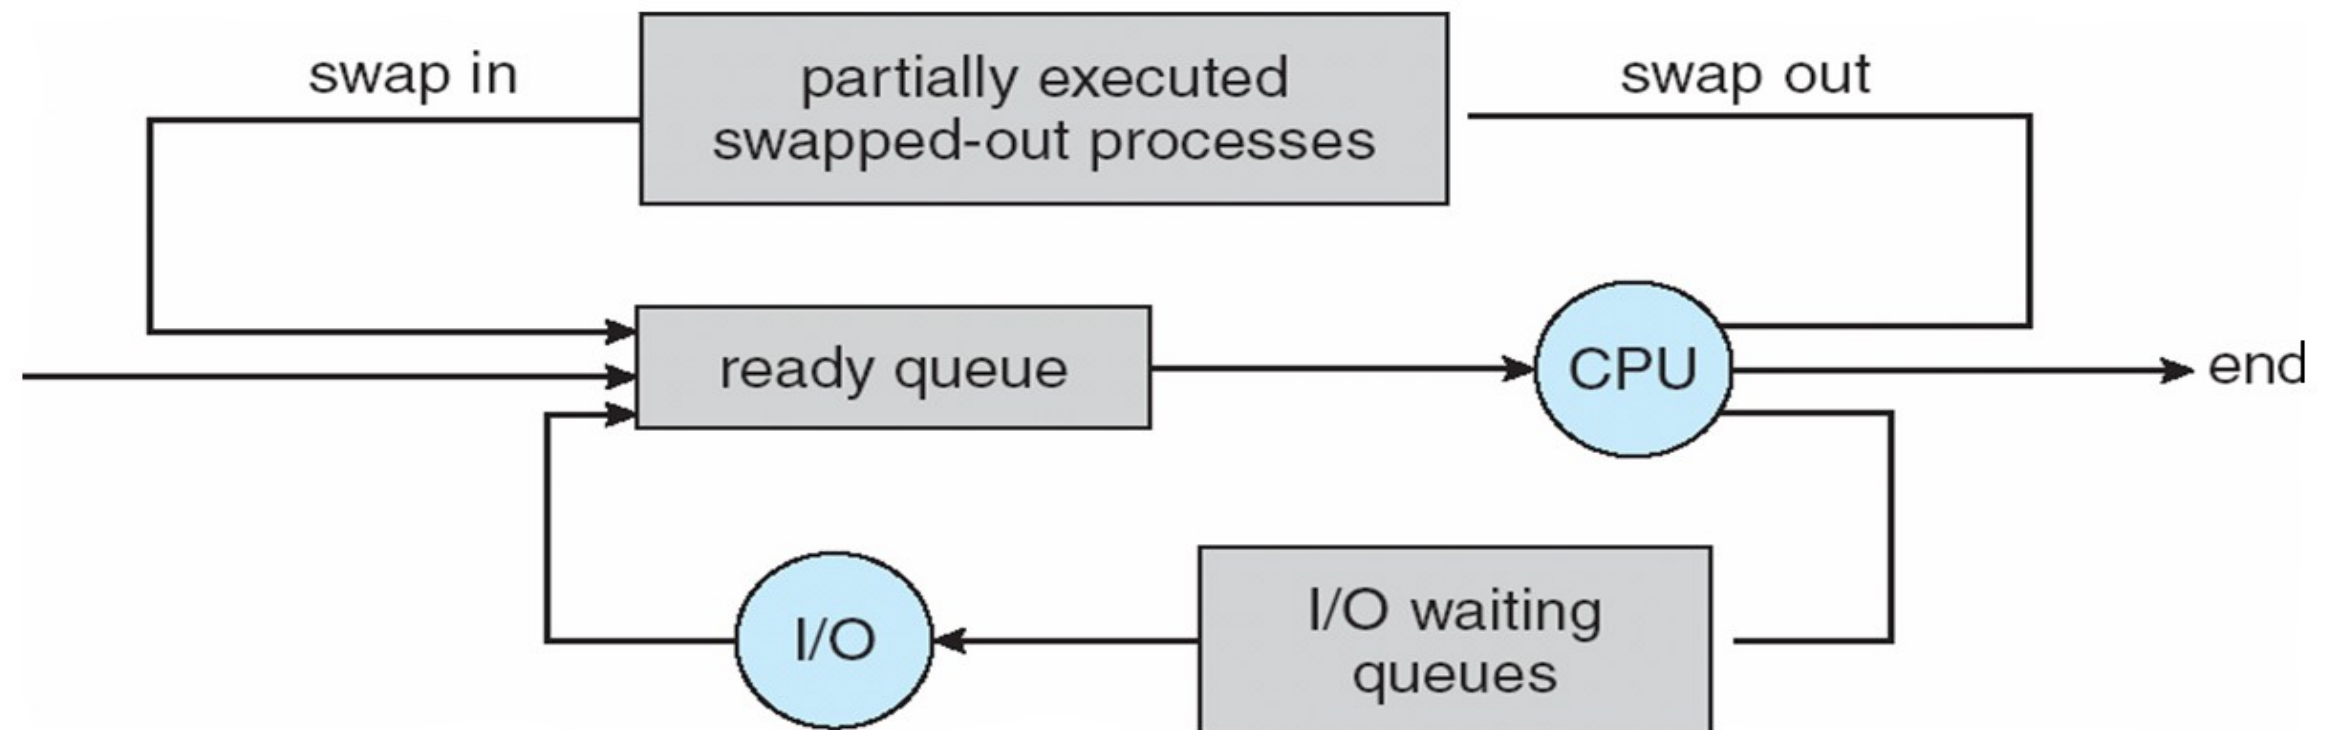
Schedulers (Cont.)


- ***Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)***
- ***Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)***
- ***The long-term scheduler controls the degree of multiprogramming***
- ***Processes can be described as either:***
 - I/O-bound process – spends more time doing I/O than computations, many short CPU bursts
 - CPU-bound process – spends more time doing computations; few very long CPU bursts





Addition of Medium Term Scheduling






Context Switch

- ***When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**.***
- *****Context** of a process represented in the PCB***
- ***Context-switch time is overhead; the system does no useful work while switching***
 - The more complex the OS and the PCB -> longer the context switch
- ***Time dependent on hardware support***
 - Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once





Process Creation

- ***Parent process create children processes, which, in turn create other processes, forming a tree of processes***
- ***Generally, process identified and managed via a process identifier (pid)***
- ***Resource sharing***
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- ***Execution***
 - Parent and children execute concurrently
 - Parent waits until children terminate





Process Creation (Cont.)

■ ***Address space***

- Child duplicate of parent
- Child has a program loaded into it

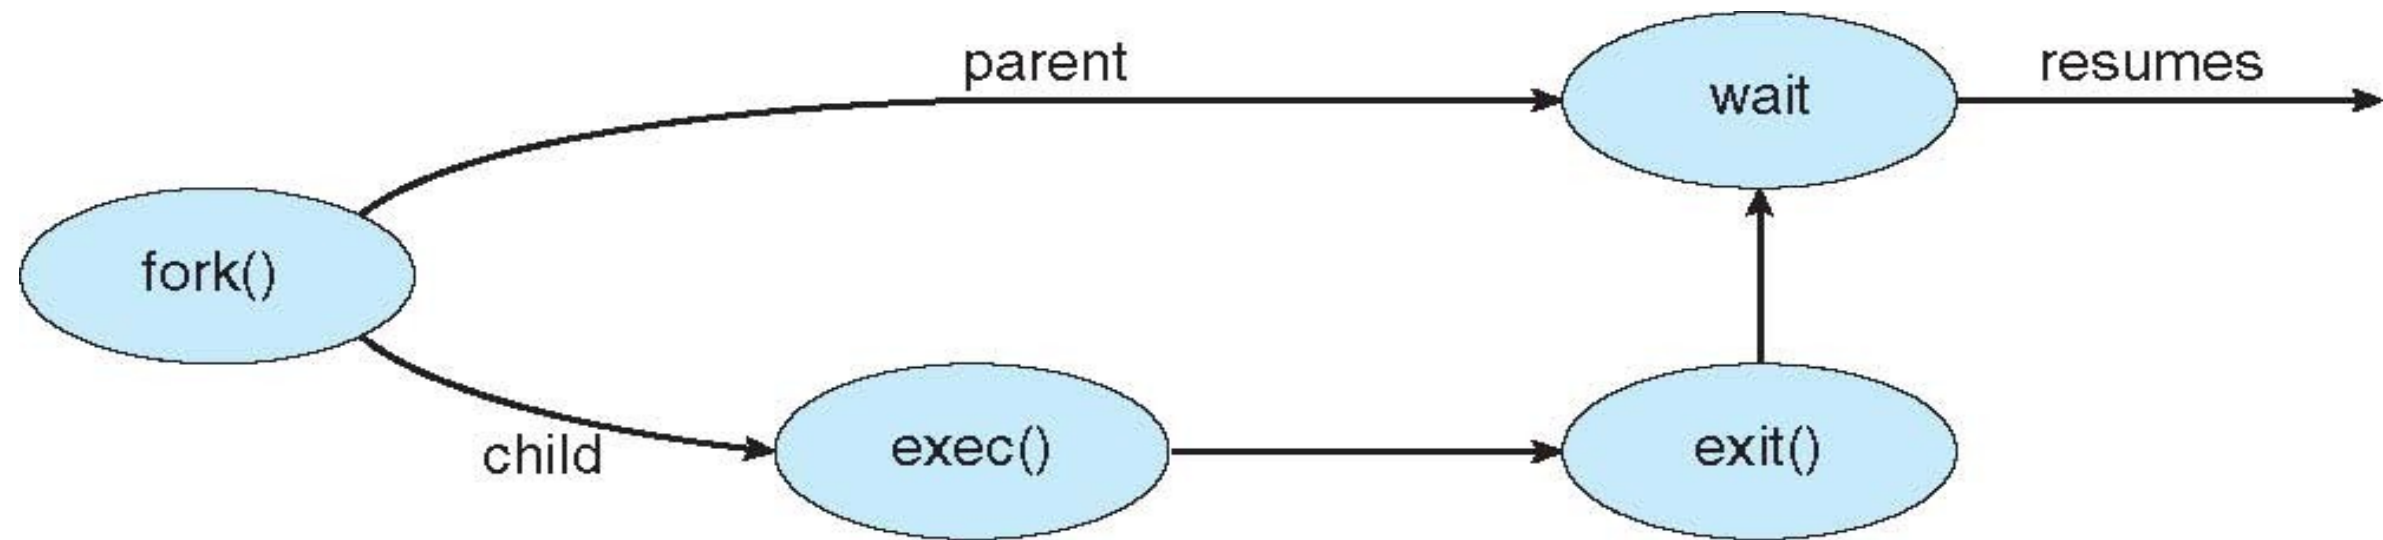
■ ***UNIX examples***

- **fork** system call creates new process
- **exec** system call used after a **fork** to replace the process' memory space with a new program





Process Creation





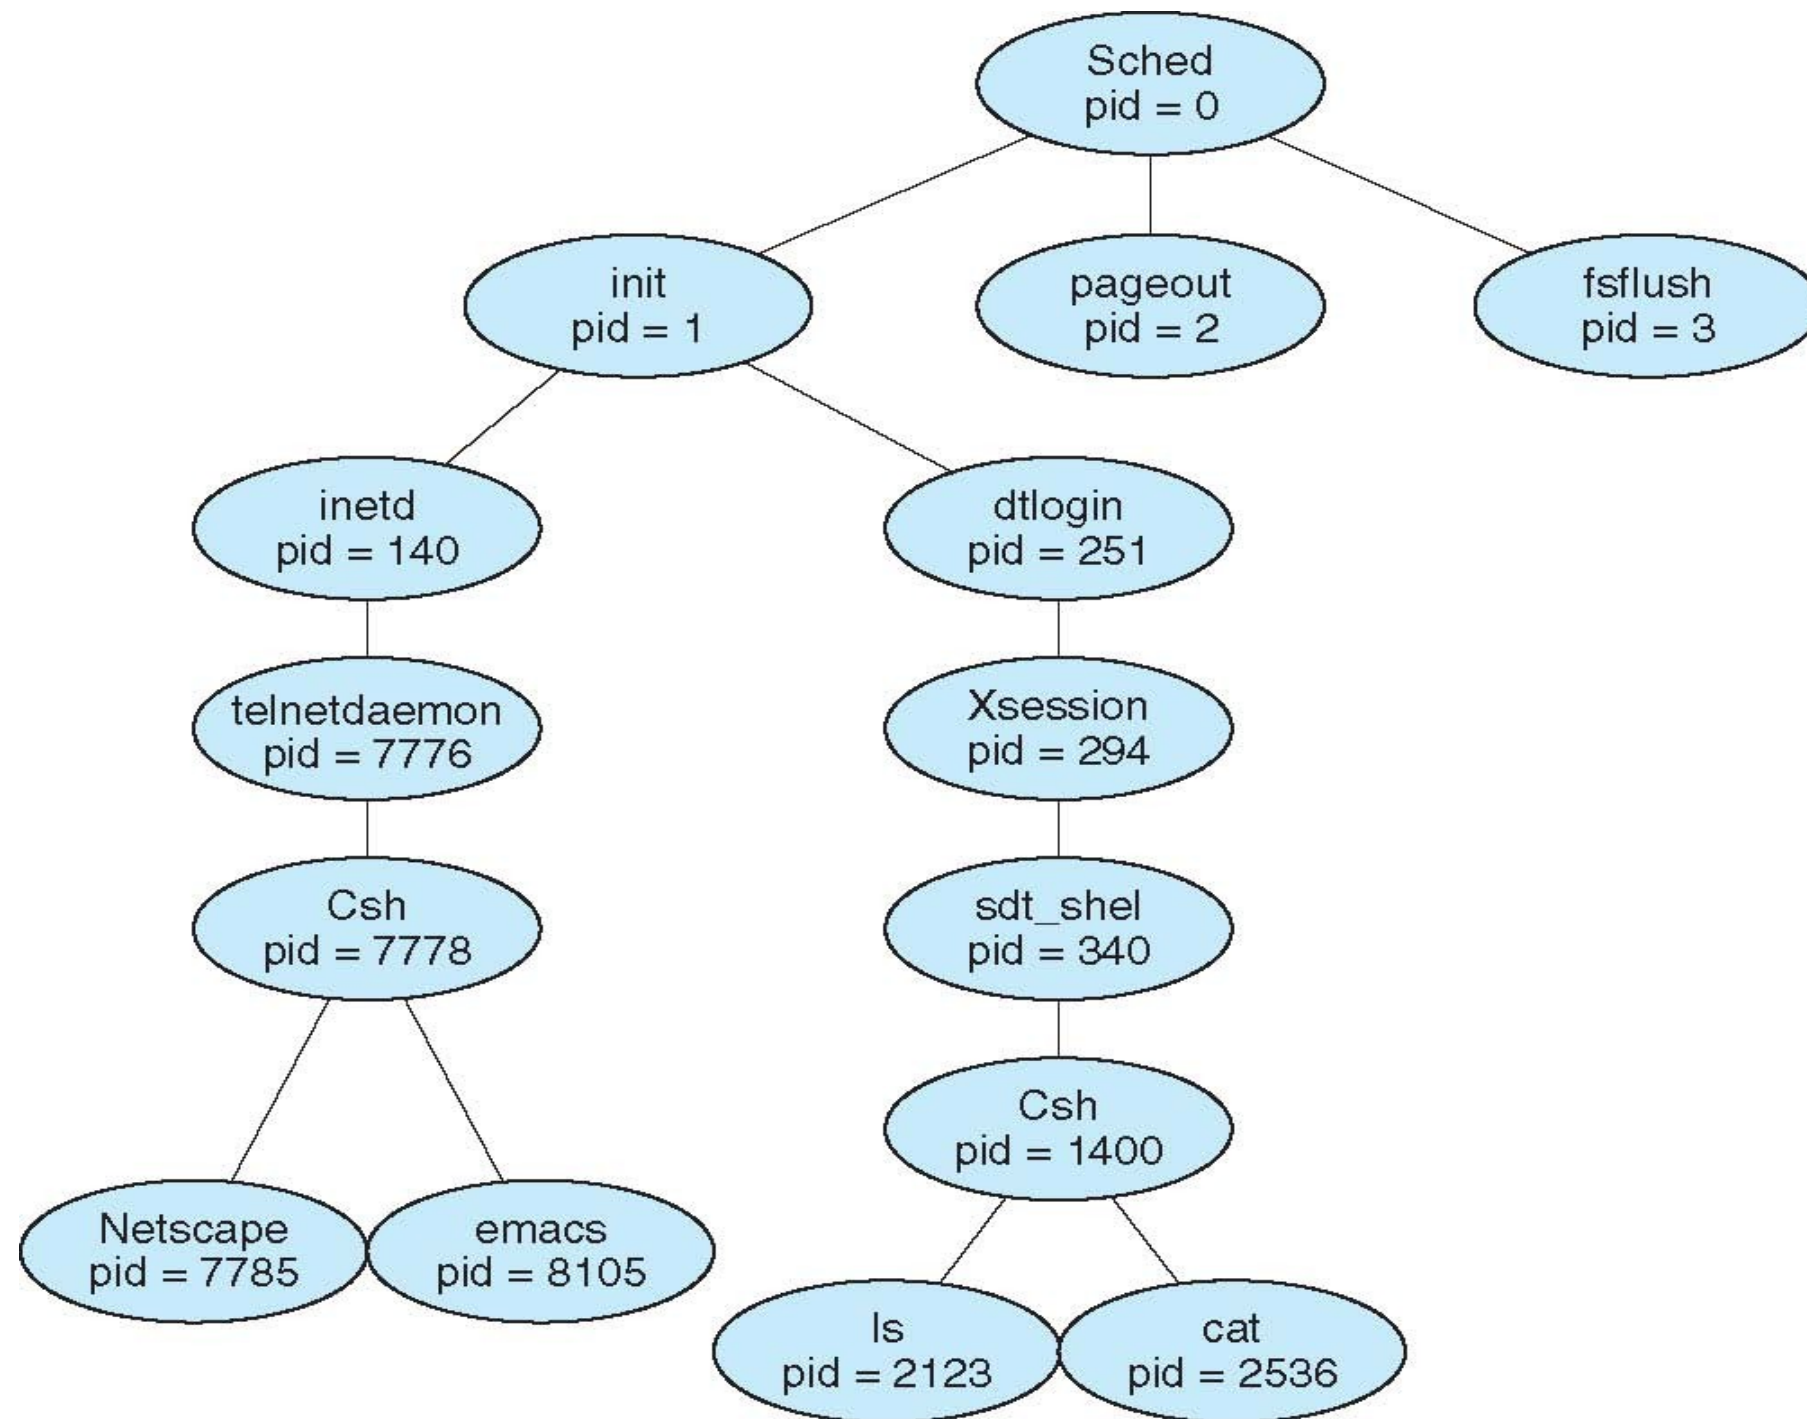
C Program Forking Separate Process

```
#include <sys/types.h>  
#include <stdio.h>  
#include <unistd.h>  
int main()  
{  
    pid_t pid;  
    /* fork another process */  
    pid = fork();  
    if (pid < 0) { /* error occurred */  
        fprintf(stderr, "Fork Failed");  
        return 1;  
    }  
    else if (pid == 0) { /* child process */  
        execlp("/bin/ls", "ls", NULL);  
    }  
    else { /* parent process */  
        /* parent will wait for the child */  
        wait (NULL);  
        printf ("Child Complete");  
    }  
    return 0;  
}
```





A Tree of Processes on Solaris





Process Termination

■ ***Process executes last statement and asks the operating system to delete it (exit)***

- Output data from child to parent (via **wait**)
- Process' resources are deallocated by operating system

■ ***Parent may terminate execution of children processes (abort)***

- Child has exceeded allocated resources
- Task assigned to child is no longer required
- If parent is exiting
 - ▶ Some operating system do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**






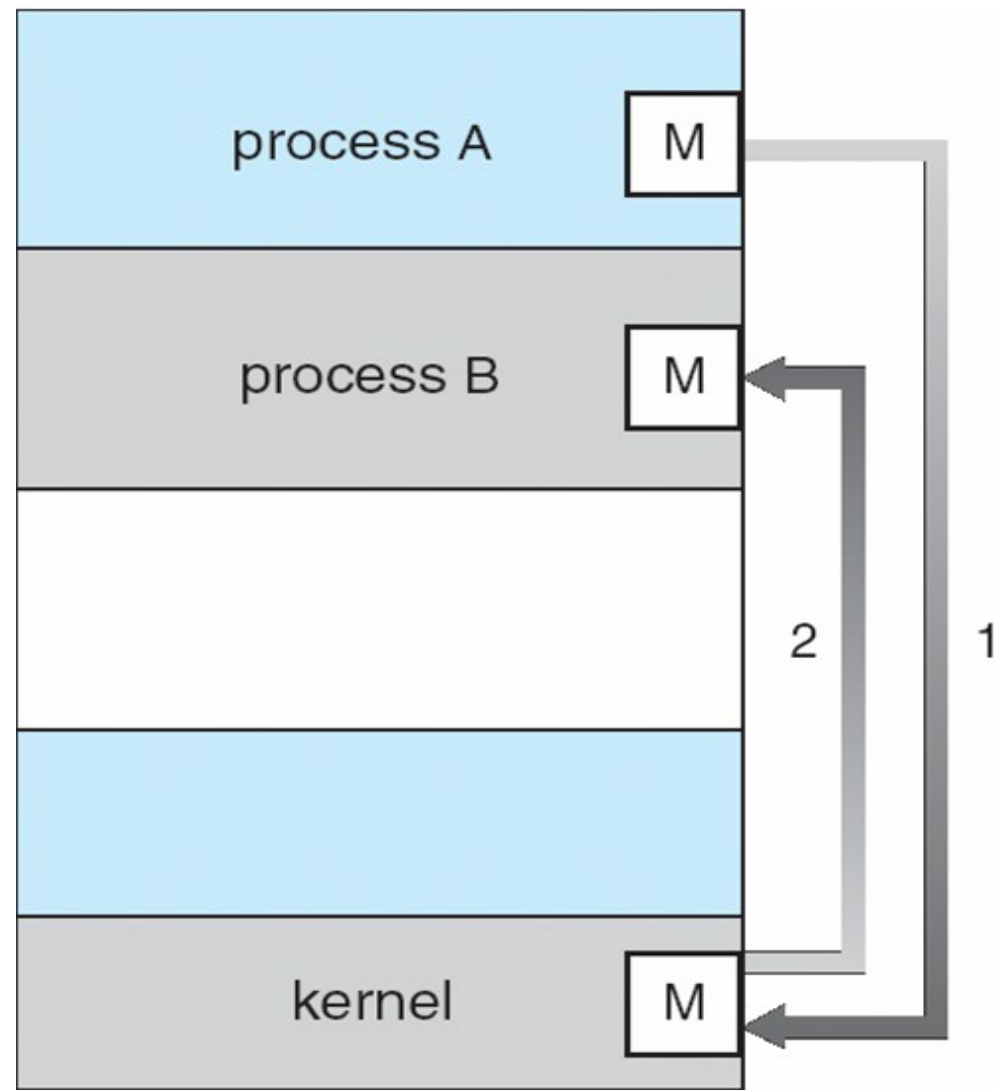
Interprocess Communication

- ***Processes within a system may be independent or cooperating***
- ***Cooperating process can affect or be affected by other processes, including sharing data***
- ***Reasons for cooperating processes:***
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- ***Cooperating processes need interprocess communication (IPC)***
- ***Two models of IPC***
 - Shared memory
 - Message passing

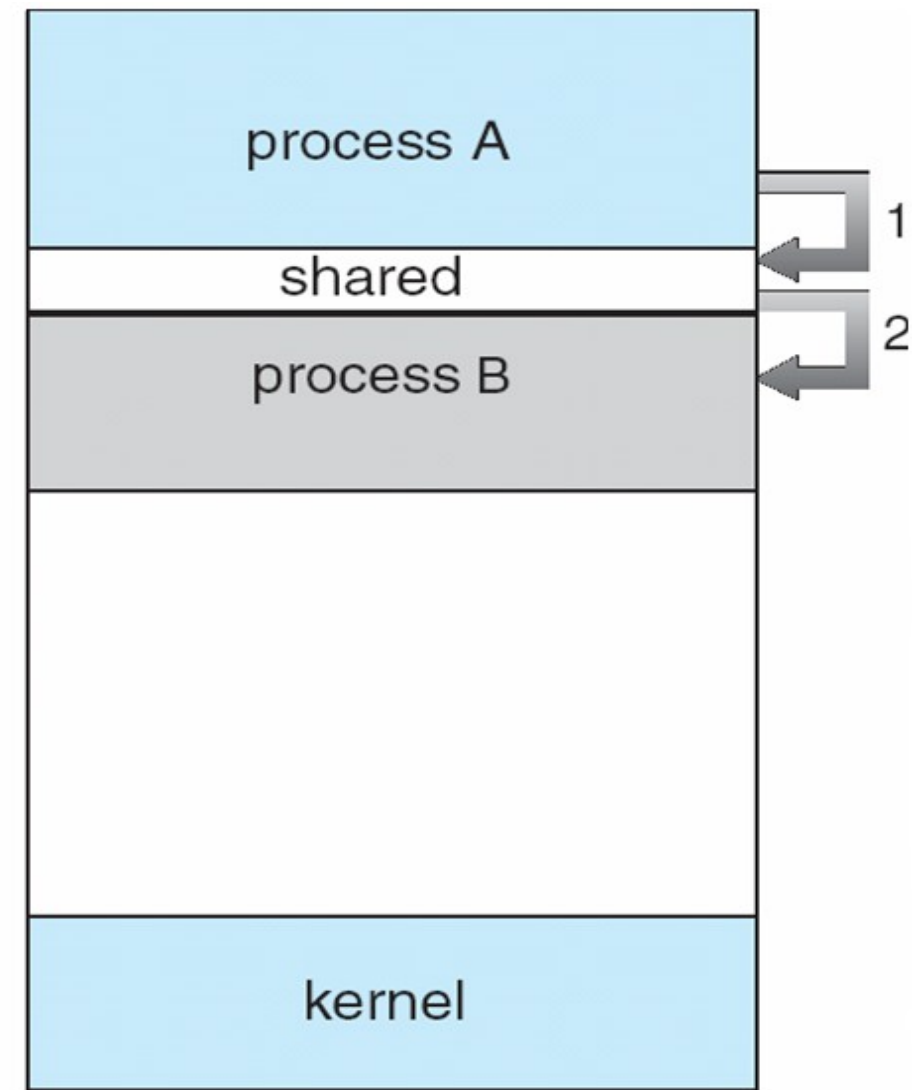




Communications Models




(a)



(b)





Cooperating Processes

- ***Independent process cannot affect or be affected by the execution of another process***
- ***Cooperating process can affect or be affected by the execution of another process***
- ***Advantages of process cooperation***
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience





Producer-Consumer Problem

■ ***Paradigm for cooperating processes, producer process produces information that is consumed by a consumer process***

- *unbounded-buffer* places no practical limit on the size of the buffer
- *bounded-buffer* assumes that there is a fixed buffer size





Bounded-Buffer – Shared-Memory Solution

■ ***Shared data***

```
#define BUFFER_SIZE 10  
typedef struct {  
    . . .  
} item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```

■ ***Solution is correct, but can only use ***BUFFER_SIZE-1*** elements***





Bounded-Buffer – Producer

```
while (true) {  
    /* Produce an item */  
    while (((in + 1) % BUFFER SIZE) == out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```





Bounded Buffer – Consumer

```
while (true) {  
    while (in == out) ;  
    // do nothing--nothing to consume  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```





Interprocess Communication – Message Passing

- ***Mechanism for processes to communicate and to synchronize their actions***
- ***Message system - processes communicate with each other without resorting to shared variables***
- ***IPC facility provides two operations:***
 - `send(message)` – message size fixed or variable
 - `receive(message)`
- ***If P and Q wish to communicate, they need to:***
 - establish a *communication link* between them
 - exchange messages via send/receive
- ***Implementation of communication link***
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)





Implementation Questions

- ***How are links established?***
- ***Can a link be associated with more than two processes?***
- ***How many links can there be between every pair of communicating processes?***
- ***What is the capacity of a link?***
- ***Is the size of a message that the link can accommodate fixed or variable?***
- ***Is a link unidirectional or bi-directional?***





Direct Communication


■ ***Processes must name each other explicitly:***

- **send** (*P, message*) – send a message to process P
- **receive**(*Q, message*) – receive a message from process Q

■ ***Properties of communication link***

- Links are established automatically
- A link is associated with exactly one pair of communicating processes
- Between each pair there exists exactly one link
- The link may be unidirectional, but is usually bi-directional





Indirect Communication

■ ***Messages are directed and received from mailboxes (also referred to as ports)***

- Each mailbox has a unique id
- Processes can communicate only if they share a mailbox

■ ***Properties of communication link***

- Link established only if processes share a common mailbox
- A link may be associated with many processes
- Each pair of processes may share several communication links
- Link may be unidirectional or bi-directional





Indirect Communication

■ **Operations**

- create a new mailbox
- send and receive messages through mailbox
- destroy a mailbox

■ **Primitives are defined as:**

send(A, message) - send a message to mailbox A

receive(A, message) - receive a message from mailbox A





Indirect Communication


■ ***Mailbox sharing***

- P_1 , P_2 , and P_3 share mailbox A
- P_1 sends; P_2 and P_3 receive
- Who gets the message?

■ ***Solutions***

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Synchronization

- ***Message passing may be either blocking or non-blocking***
- ***Blocking is considered synchronous***
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- ***Non-blocking is considered asynchronous***
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null





Buffering

■ ***Queue of messages attached to the link; implemented in one of three ways***

1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
2. Bounded capacity – finite length of n messages
Sender must wait if link full
3. Unbounded capacity – infinite length
Sender never waits





Examples of IPC Systems - POSIX

■ ***POSIX Shared Memory***

- Process first creates shared memory segment
`segment_id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
- Process wanting access to that shared memory must attach to it
`shared_memory = (char *) shmat(id, NULL, 0);`
- Now the process could write to the shared memory
`sprintf(shared_memory, "Writing to shared memory");`
- When done a process can detach the shared memory from its address space
`shmdt(shared_memory);`





Examples of IPC Systems - Mach

■ ***Mach communication is message based***

- Even system calls are messages
- Each task gets two mailboxes at creation- Kernel and Notify
- Only three system calls needed for message transfer

`msg_send()`, `msg_receive()`, `msg_rpc()`

- Mailboxes needed for communication, created via
`port_allocate()`





Examples of IPC Systems – Windows XP

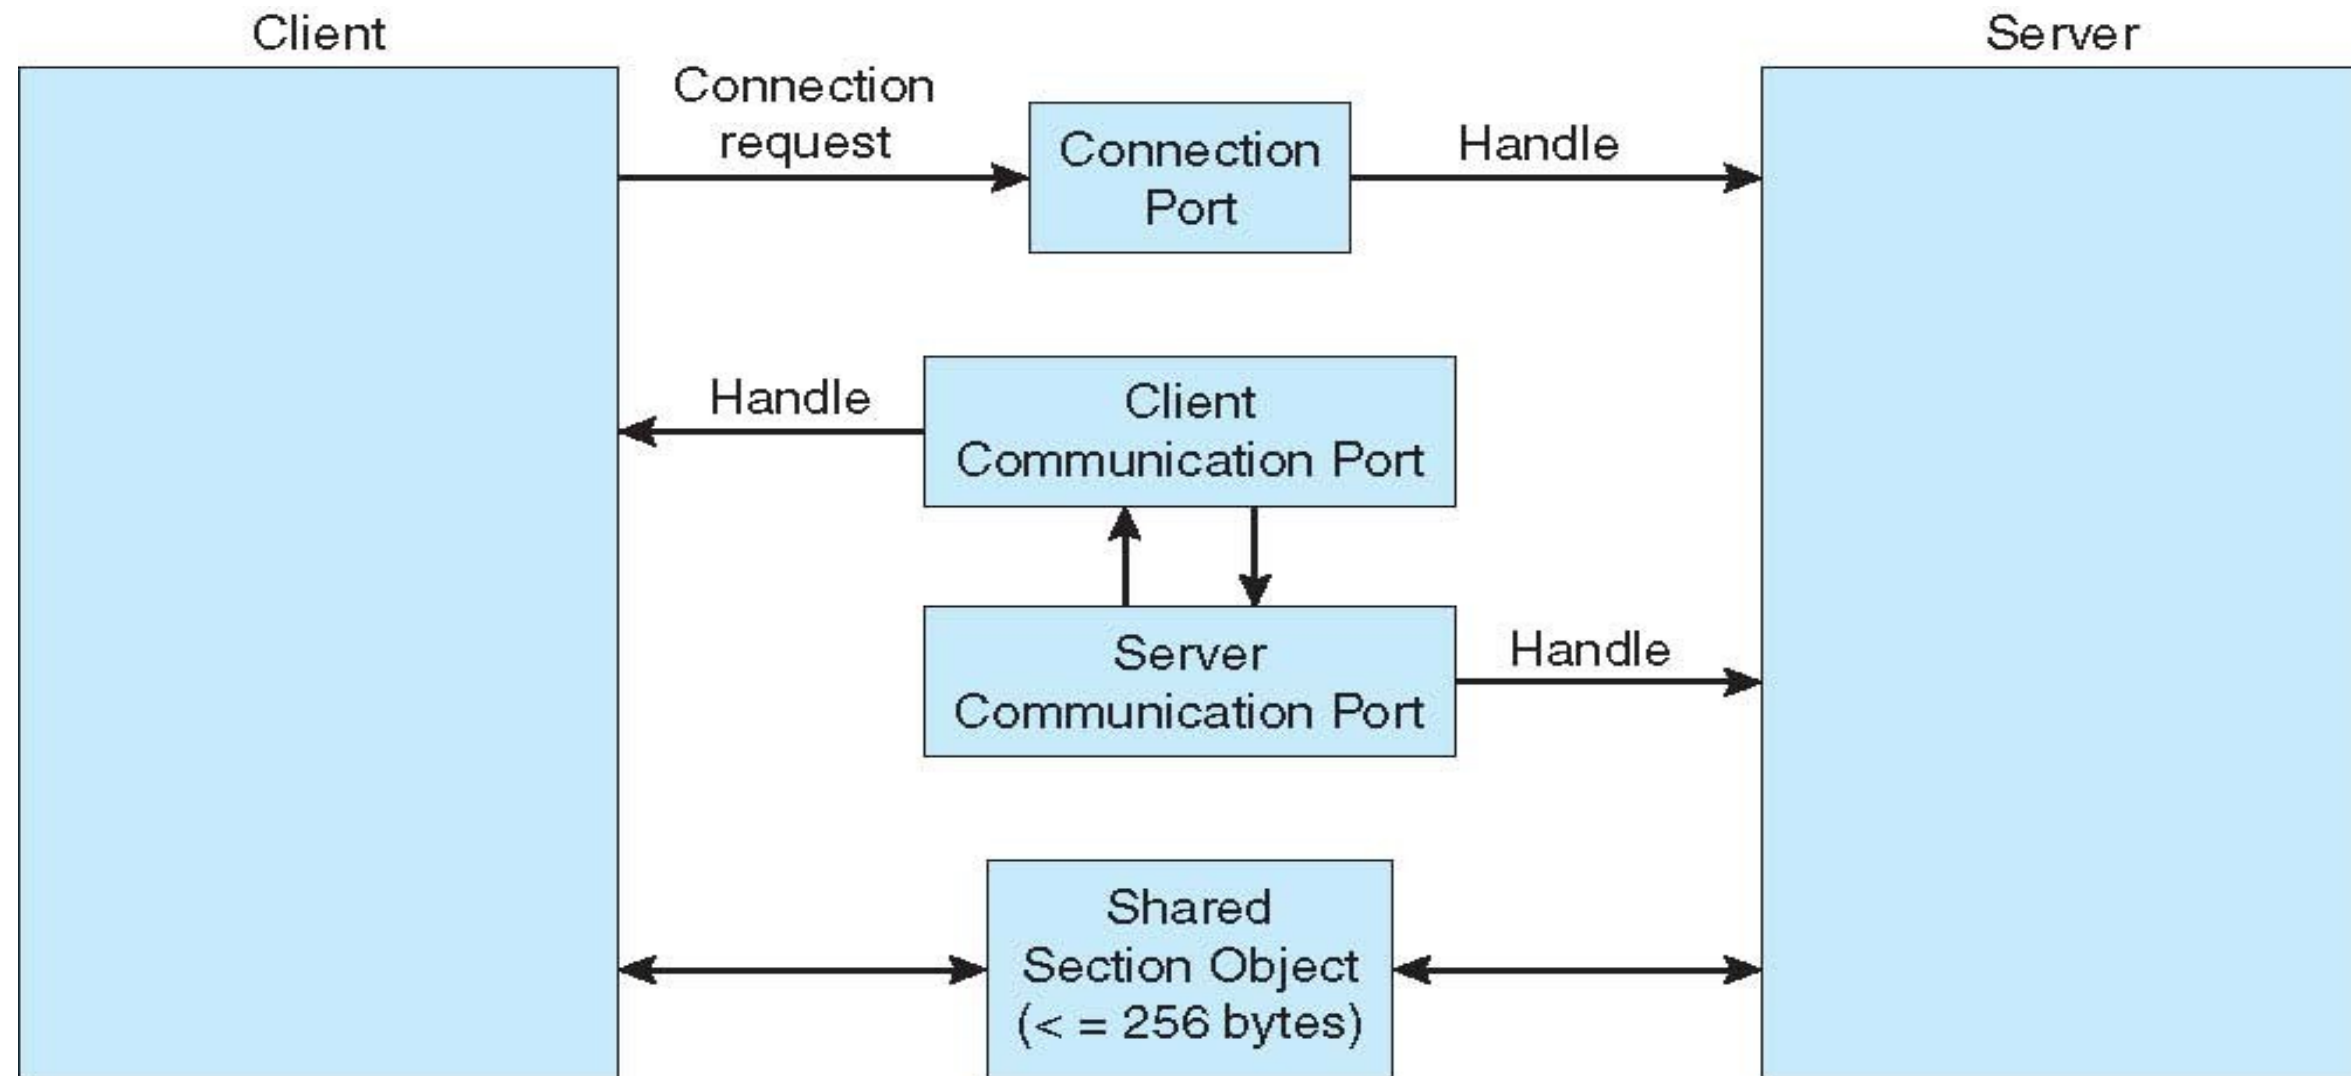
■ ***Message-passing centric via **local procedure call (LPC)** facility***

- Only works between processes on the same system
- Uses ports (like mailboxes) to establish and maintain communication channels
- Communication works as follows:
 - ▶ The client opens a handle to the subsystem's connection port object.
 - ▶ The client sends a connection request.
 - ▶ The server creates two private communication ports and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Local Procedure Calls in Windows XP





Communications in Client-Server Systems

- ***Sockets***
- ***Remote Procedure Calls***
- ***Pipes***
- ***Remote Method Invocation (Java)***





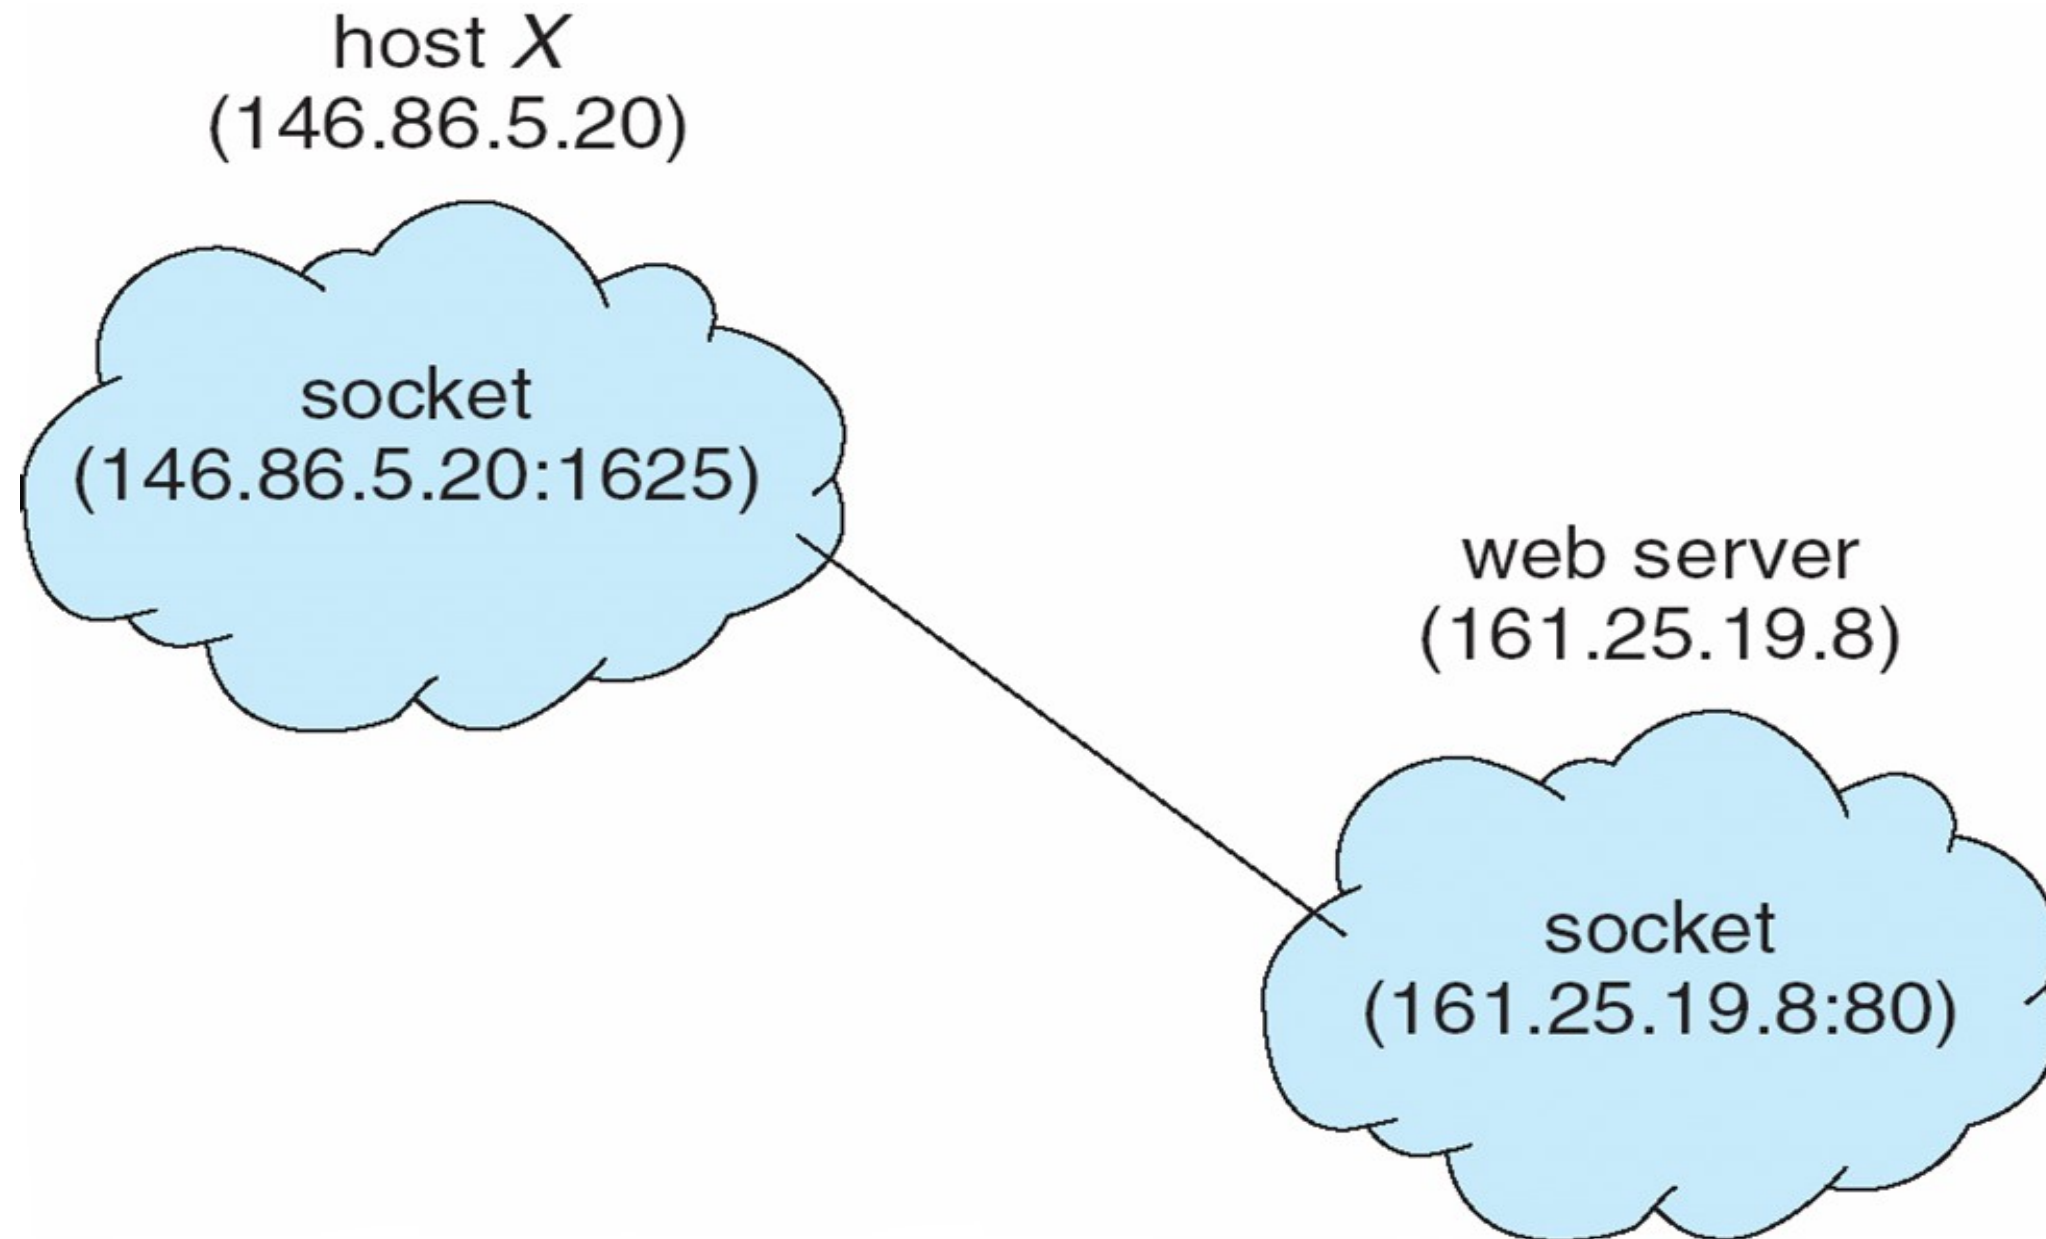
Sockets

- ***A **socket** is defined as an endpoint for communication***
- ***Concatenation of IP address and port***
- ***The socket 161.25.19.8:1625 refers to port 1625 on host 161.25.19.8***
- ***Communication consists between a pair of sockets***





Socket Communication





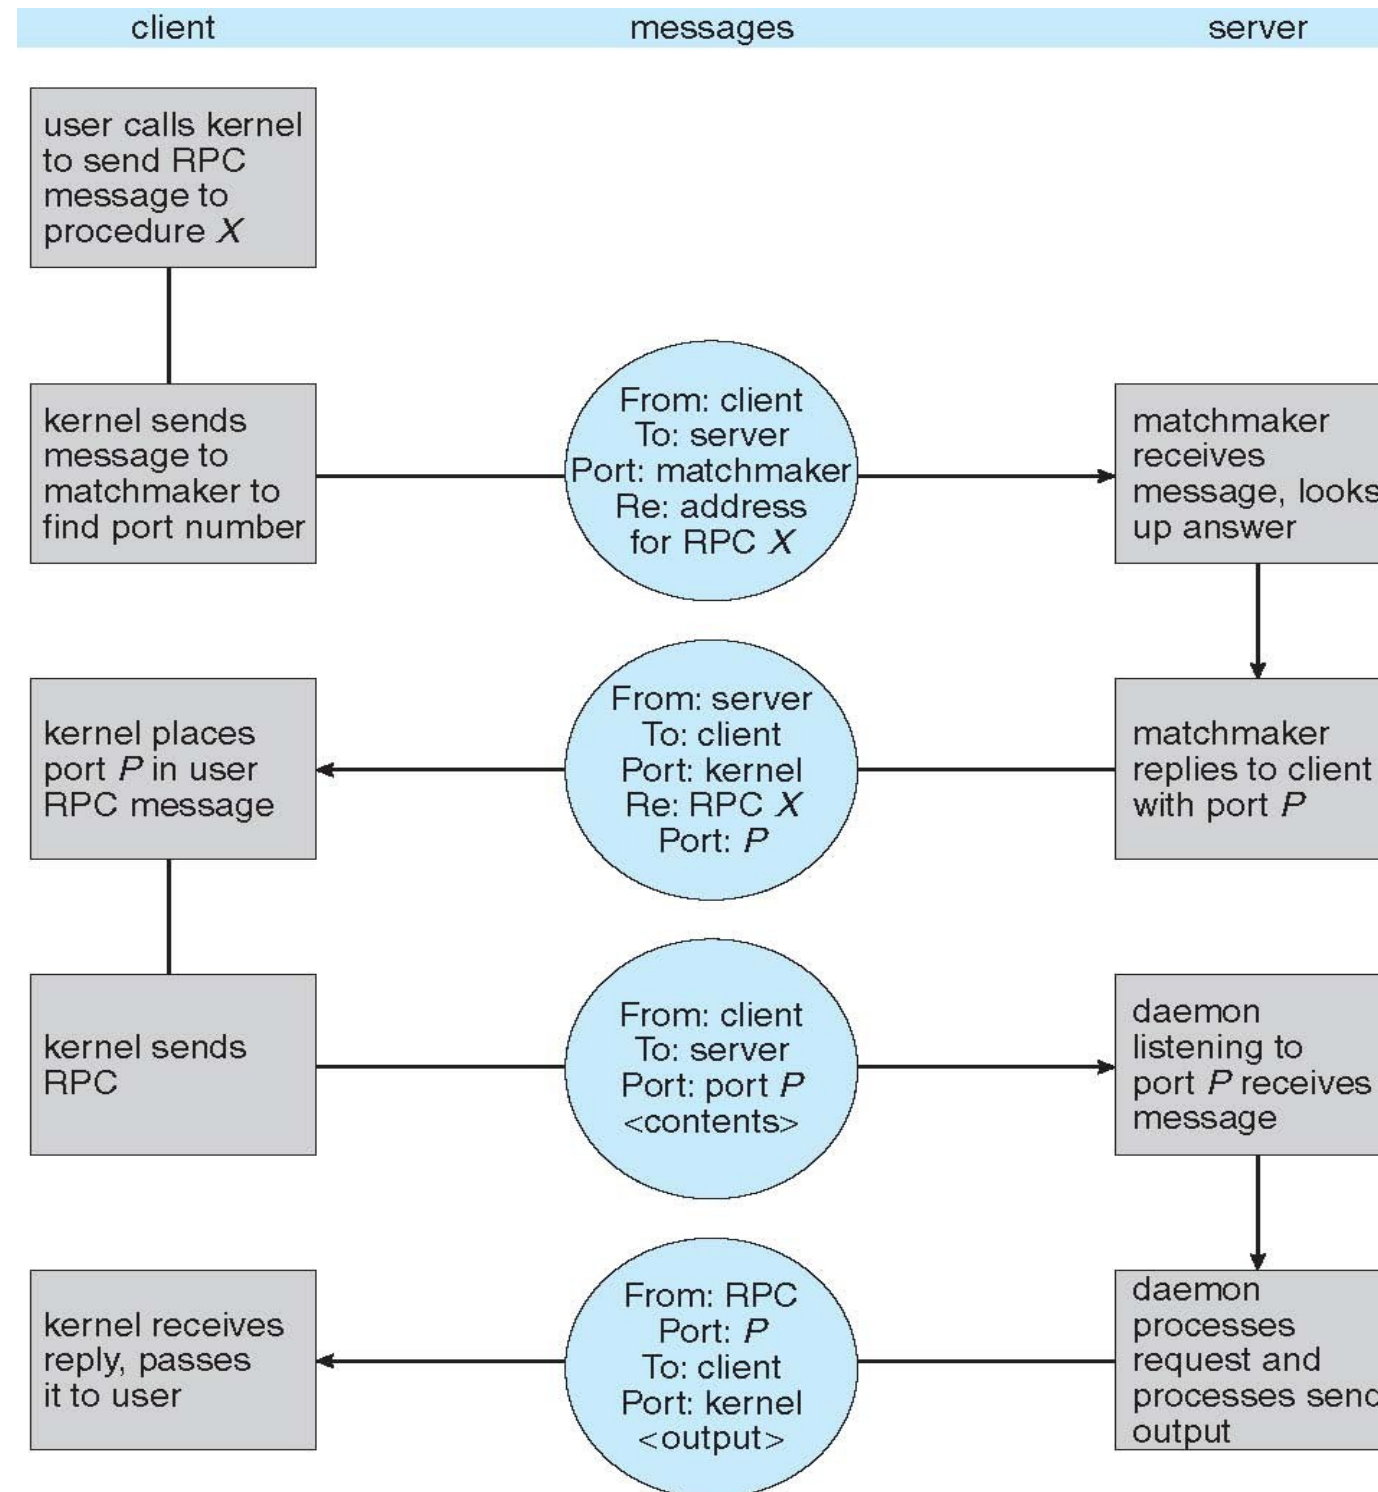
Remote Procedure Calls

- ***Remote procedure call (RPC) abstracts procedure calls between processes on networked systems***
- ***Stubs - client-side proxy for the actual procedure on the server***
- ***The client-side stub locates the server and marshalls the parameters***
- ***The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server***





Execution of RPC






Pipes

■ ***Acts as a conduit allowing two processes to communicate***

■ ***Issues***

- Is communication unidirectional or bidirectional?
- In the case of two-way communication, is it half or full-duplex?
- Must there exist a relationship (i.e. parent-child) between the communicating processes?
- Can the pipes be used over a network?






Ordinary Pipes

- ***Ordinary Pipes allow communication in standard producer-consumer style***
- ***Producer writes to one end (the write-end of the pipe)***
- ***Consumer reads from the other end (the read-end of the pipe)***
- ***Ordinary pipes are therefore unidirectional***
- ***Require parent-child relationship between communicating processes***





Ordinary Pipes





Named Pipes

- ***Named Pipes are more powerful than ordinary pipes***
- ***Communication is bidirectional***
- ***No parent-child relationship is necessary between the communicating processes***
- ***Several processes can use the named pipe for communication***
- ***Provided on both UNIX and Windows systems***



End of Chapter 3

